

Link to repo = <https://github.com/scandum/wolfsort>

The workload I chose is a benchmark of a sorting algorithm that the owner of the repository had made. The sorting algorithm in mind is called **wolfsort**, it is a hybrid sorting algorithm that combines properties of radix sort, quicksort and merge sort. The workload runs a benchmark that compares wolfsort to other hybrid sorting algorithms.

Downloading and running the workload

To download the repository, run the following command

```
git clone https://github.com/scandum/wolfsort
```

After cloning the repository, you can run and compile the benchmark as such:

```
cd wolfsort/src
gcc -O3 bench.c
./a.out
```

To compile it into a RISC-V binary run the following compile line

```
riscv64-linux-gnu-gcc -O3 -static bench.c -o bench.rv
```

The bash script I used to run the workload is the following code

```
ESESC_BIN=${1:-../main/esesc}

export ESESC_ReportFile="part2Report"
export ESESC_BenchName="./wolfsort/src/bench.rv"
if [ -f $ESESC_BIN ]; then
    $ESESC_BIN
else
    $ESESC_BenchName
fi
exit 0
```

Initial Results

The initial run of the benchmark showed an IPC of 0.39, and a total instruction count of 362,993,011 instructions, with a running time of 676 seconds. The full report given by esesc can be seen below

```
*****
# File : esesc_part3Report1.xT6px0 : Wed Nov 30 16:22:46 2022
*****
Sampler 0 (Procs 0)
*****
Rabbit Marcup Detail Timing Total KIPS
KIPS 94029 N/A 251972 606 10775
Time 5.9% 0.0% 1.8% 92.2% : Sim Time (s) 676.409 Exe 551.860 ms Sim (1700MHz)
Inst 51.8% 0.0% 43.0% 5.2% : Approx Total Time 10642.763 ms Sim (1700MHz)
*****
Proc : Delay : Avg.Time : BPType : Total : RAS : BPred : BTB : iBTB : BTAC : WasteRatio : MPKI
0 : 3 : 12.809 : 2bit : 89.73% : 99.99% of 8.36% : 88.28% of 87.64% : 88.15% of 49.47% : 0.00% : 0.00% : 0.00% : 23.50
0 : 4 : 12.809 : 2level : 89.73% : 0.00% of 0.00% : 88.69% of 90.80% : 69.19% of 0.09% : 49.38% : ( 3.22% fixed) : 23.50
*****
Proc : nCommIt : nInst : AALU : BALU : CALU : LALU : SALU : LD Fwd : Replay : Worst Unit (clk)
0 : 362993011 : 362993052 : 47.87% : 24.80% : 0.05% : 14.37% : 12.91% : 0.02% : N/A : 0.00
*****
Proc IPC uIPC Active Cycles Busy LDQ STQ IWin ROB Regs IO maxBr MlsBr Br4Clk brDelay
0 0.39 0.39 1.00 938162540 19.3 0.0 51.5 0.2 0.0 1.0 0.0 0.0 0.0 0.0 0.0 8.8
*****
Cache Occ AvgMenLat MenAccesses MissRate ( RD , WR , BUS)
IL1(0) 0.0 2.0 262853301 0.0% 0.0% (100.0%, 0.0%, 0.0%) 29.8 0.0 GB/s
*****
DL1(0) 0.0 98.7 99024207 14.0% 53.7% ( 77.7%, 44.3%, 0.0%) 11.2 0.0 GB/s
*****
L2(0) 0.0 186.3 13904016 92.4% 92.4% ( 7.6%, 0.0%, 0.0%) 1.6 0.0 GB/s
Memory(0) 0.0 60.0 13220519 0.0% 0.0% (100.0%,100.0%, 0.0%) 1.5 0.0 GB/s
*****
```

Finding the bottleneck

By looking at the report, it was clear that the L2 cache was underperforming heavily with a 90% miss rate on both types of misses. I double-checked that this was the case by testing changes with the L1 cache since it also had a relatively high miss rate at 50%. Some of the changes I tried were increasing both block size and associativity of the L1 cache, these changes resulted in little to no speedup.